

Ay190 – Worksheet 9
Daniel DeFelippis
Date: February 11, 2014

Solving Large Systems of Linear Equations

1

Writing a script to load these data files as matrix objects is fairly easy using `np.loadtxt` as shown below.

```
LSE1_m = np.loadtxt(' ./LSE1_m.dat')
```

Typing "`LSE1_m.shape`" gives the dimensions of the matrix `LSE1_m`. Doing this, we see that the matrices `LSEi_m` with $i = 1, \dots, 5$ are all square matrices with number of rows (=number of columns) being 10, 100, 200, 1000, and 2000 respectively. Using the function "`slogdet`" located in NumPy's `linalg` module, we can calculate the natural log of the determinant to make sure that the determinant isn't 0 so the LSE is solvable. This is indeed true.

2

I wrote my own code that implements the Gauss algorithm. I defined two functions, one which gets the LSE into the desired triangular form, and the other which backsubstitutes to find all of the x_i in the LSE equation $Ax = \mathbf{b}$. I ran the algorithm on the contrived example

$$\begin{pmatrix} 2 & 5 & -3 \\ 1 & -7 & 4 \\ 6 & -2 & 2 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ x_3 \end{pmatrix} = \begin{pmatrix} 3 \\ -1 \\ 4 \end{pmatrix}$$

with the known solution of

$$\mathbf{x} = \begin{pmatrix} 1 \\ 2 \\ 3 \end{pmatrix}$$

to see if the algorithm returned the correct answer. It did!

So, I could then use it on the much larger LSE matrices. I timed how long running the algorithm on each LSE took using the "`time`" function in the `time` module to store the starting and stopping time and then print the difference between the two (the times are measured in seconds). The timing results are given for a particular run of the algorithm on the five matrices.

LSE	Size	Time (seconds)
1	(10, 10)	0.00103902816772
2	(100, 100)	0.173527002335
3	(200, 200)	0.333551883698
4	(1000, 1000)	10.1985230446
5	(2000, 2000)	50.3575429916

The actual values did sometimes vary, probably just due to how many other random processes happened to be running on my laptop at that time.

3

Next, I try doing solving the same systems with NumPy's "solve" function in its linalg module. I get faster results compared to Gaussian Elimination as shown in the table below.

LSE	Size	Time (Gauss)	Time (NumPy)	Ratio of Times
1	(10, 10)	0.00103902816772	0.000102043151855	10.182242990655809
2	(100, 100)	0.173527002335	0.00051212310791	338.8384543770586
3	(200, 200)	0.333551883698	0.0143749713898	23.203655482380803
4	(1000, 1000)	10.1985230446	0.521645069122	19.55069384968118
5	(2000, 2000)	50.3575429916	3.08289694786	16.334487932383144

NumPy's solver is clearly much better. For all sizes, it is at least an order of magnitude faster.

SciPy also has a bunch of solvers in its "sparse.linalg" module. However, they are all iterative, and the ones I tried (spsolve and cg) were nowhere near NumPy's solver's speed. I also tried SciPy's "solve" function located in its own linalg module, and it performed almost identically well as (but no better than) NumPy's solver.